


Question to be answered (Q2.4)

How do D-loss and G-loss change during training? Visualize how the D-loss and D-loss change during training and explain why.

GANs try to replicate a probability distribution. They should therefore use loss functions that reflect the distance between the distribution of the data generated by the GAN and the distribution of the real data. In the paper that introduced GANs, the generator tries to minimize the following function while the discriminator tries to maximize it during training.

As we see in the picture below the visualisation of the D-loss and the G-loss.



While the discriminator is trained, it classifies both the real data and the fake data from the generator. It penalizes itself for misclassifying a real instance as fake, or a fake instance (created by the generator) as real, by maximizing the below function.

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right]$$

$\log(D(x))$ refers to the probability that the generator is rightly classifying the real image,

maximizing $\log(1-D(G(z)))$ would help it to correctly label the fake image that comes from the generator.

While the generator is trained, it samples random noise and produces an output from that noise. The output then goes through the discriminator and gets classified as either “Real” or “Fake” based on the ability of the discriminator to tell one from the other. The generator loss is then calculated from the discriminator’s classification – it gets rewarded if it successfully fools the discriminator, and gets penalized otherwise.

The following equation is minimized to training the generator:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)})))$$

Exercise to be conducted (E1.4)

To get and show the low-resolution image

```
import os
```

```
import pickle
```

```
import time
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
import torch
```

```
import torch.nn as nn
```

```
import torchvision.datasets as datasets
```

```
import torchvision.transforms as transforms
```

```
from torch.utils.data import DataLoader
```

```
from tqdm import tqdm
```

```
class Generator(nn.Module):
```

```
    """Image generator
```

```
    Takes a noise vector as input and synthesises a single channel image accordingly
```

```
    """
```

```
    def __init__(self, input_dims, output_dims):
```

```
        """Init function
```

```
        Declare the network structure as indicated in CW2 Guidance
```

```
        Arguments:
```

```
            input_dims {int} -- Dimension of input noise vector
```

```
            output_dims {int} -- Dimension of the output vector (flatten image)
```

```
        """
```

```
        super(Generator, self).__init__()
```

```

    """ TODO: Change the architecture and value as CW2 Guidance required

    self.fc0 = nn.Sequential(
        nn.Linear(input_dims, 128),
        nn.LeakyReLU(0.2))

    # output hidden layer
    self.fc1 = nn.Sequential(
        nn.Linear(128, output_dims),
        nn.Tanh())

def forward(self, x):
    """Forward function

    Arguments:
        x {Tensor} -- a batch of noise vectors in shape (<batch_size>x<input_dims>)

    Returns:
        Tensor -- a batch of flatten image in shape (<batch_size>x<output_dims>)
    """

    """ TODO: modify to be consistent with the network structure

    x = self.fc0(x)
    x = self.fc1(x)

    return x

class Discriminator(nn.Module):
    """Image discriminator

    Takes a image as input and predict if it is real from the dataset or fake synthesised by the
    generator
    """

    def __init__(self, input_dims, output_dims=1):

```

```
"""Init function
```

Declare the discriminator network structure as indicated in CW2 Guidance

Arguments:

input_dims {int} -- Dimension of the flatten input images

Keyword Arguments:

output_dims {int} -- Predicted probability (default: {1})

```
"""
```

```
super(Discriminator, self).__init__()
```

```
### TODO: Change the architecture and value as CW2 Guidance required
```

```
self.fc0 = nn.Sequential(  
    nn.Linear(input_dims, 784),  
    nn.LeakyReLU(0.2),  
    nn.Dropout(0.3)  
)
```

```
self.fc1 = nn.Sequential(  
    nn.Linear(784, 1),  
    nn.Sigmoid()  
)
```

```
def forward(self, x):
```

```
    """Forward function
```

Arguments:

x {Tensor} -- a batch of 2D image in shape (<batch_size>xHxW)

Returns:

Tensor -- predicted probabilities (<batch_size>)

```

"""

### TODO: modify to be consistent with the network structure

x = self.fc0(x)
x = self.fc1(x)
return x

def show_result(G_net, z_, num_epoch, show=False, save=False, path='result.png'):
    """Result visualisation

    Show and save the generated figures in the grid fashion

    Arguments:
        G_net {[nn.Module]} -- The generator instant
        z_ {[Tensor]} -- Input noise vectors
        num_epoch {[int]} -- Indicate how many epoch has the generator been trained

    Keyword Arguments:
        show {bool} -- If to display the images (default: {False})
        save {bool} -- If to store the images (default: {False})
        path {str} -- path to store the images (default: {'result.png'})
    """

    ### TODO: complete the rest of part
    # hint: use plt.subplots to construct grid
    # hint: use plt.imshow and plt.savefig to display and store the images

def show_train_hist(hist, show=False, save=False, path='Train_hist.png'):
    """Loss tracker

```

Plot the losses of generator and discriminator independently to see the trend

Arguments:

hist {[dict]} -- Tracking variables

Keyword Arguments:

show {bool} -- If to display the figure (default: {False})

save {bool} -- If to store the figure (default: {False})

path {str} -- path to store the figure (default: {'Train_hist.png'})

"""

```
x = range(len(hist['D_losses']))
```

```
y1 = hist['D_losses']
```

```
y2 = hist['G_losses']
```

```
plt.plot(x, y1, label='D_loss')
```

```
plt.plot(x, y2, label='G_loss')
```

```
plt.xlabel('Epoch')
```

```
plt.ylabel('Loss')
```

```
plt.legend(loc=4)
```

```
plt.grid(True)
```

```
plt.tight_layout()
```

```
if save:
```

```
    plt.savefig(path)
```

```
if show:
```

```
    plt.show()
```

```
else:
```

```
    plt.close()
```

```
def create_noise(num, dim):
```

```
    """Noise constructor
```

```
    returns a tensor filled with random numbers from a standard normal distribution
```

```
Arguments:
```

```
    num {int} -- Number of vectors
```

```
    dim {int} -- Dimension of vectors
```

```
Returns:
```

```
    [Tensor] -- the generated noise vector batch
```

```
    """
```

```
    return torch.randn(num, dim)
```

```
if __name__ == '__main__':
```

```
    # initialise the device for training, if gpu is available, device = 'cuda', else: device = 'cpu'
```

```
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
    data_dir = './MNIST_data/'
```

```
    save_dir = './MNIST_GAN_results/'
```

```
    image_save_dir = './MNIST_GAN_results/results'
```

```
    # create folder if not exist
```

```
    if not os.path.exists(save_dir):
```

```
        os.mkdir(save_dir)
```

```
    if not os.path.exists(image_save_dir):
```

```
        os.mkdir(image_save_dir)
```

```

# training parameters : change here
batch_size = 100
learning_rate = 0.0002
epochs = 200

# parameters for Models
image_size = 28
G_input_dim = 100
G_output_dim = image_size * image_size
D_input_dim = image_size * image_size
D_output_dim = 1
hidden_size = 32

# construct the dataset and data loader
transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize(mean=(0.5,),
std=(0.5,))])
train_data = datasets.MNIST(root=data_dir, train=True, transform=transform, download=True)
train_loader = DataLoader(dataset=train_data, batch_size=batch_size, shuffle=True)

# declare the generator and discriminator networks
G_net = Generator(G_input_dim, G_output_dim).to(device)
D_net = Discriminator(D_input_dim, D_output_dim).to(device)

# Binary Cross Entropy Loss function
criterion = nn.BCELoss().to(device)

# Initialise the Optimizers
G_optimizer = torch.optim.Adam(G_net.parameters(), lr=learning_rate)
D_optimizer = torch.optim.Adam(D_net.parameters(), lr=learning_rate)

```



```

# tracking variables
train_hist = {}
train_hist['D_losses'] = []
train_hist['G_losses'] = []
train_hist['per_epoch_ptimes'] = []
train_hist['total_ptime'] = []

start_time = time.time()

# training loop
for epoch in range(epochs):
    G_net.train()
    D_net.train()
    Loss_G = []
    Loss_D = []
    epoch_start_time = time.time()
    for (image, _) in tqdm(train_loader):
        image = image.to(device)
        b_size = len(image)
        # creat real and fake labels
        real_label = torch.ones(b_size, 1).to(device)
        fake_label = torch.zeros(b_size, 1).to(device)

        # generate fake images
        data_fake = G_net(create_noise(b_size, G_input_dim).to(device))
        data_real = image.view(b_size, D_input_dim)

        # -----train the discriminator network-----
        # compute the loss for real and fake images
        output_real = D_net(data_real)
        output_fake = D_net(data_fake)
        loss_real = criterion(output_real, real_label)

```

```

loss_fake = criterion(output_fake, fake_label)
loss_d = loss_real + loss_fake

# back propagation
D_optimizer.zero_grad()
loss_d.backward()
D_optimizer.step()

# ----- train the generator network-----
data_fake = G_net(create_noise(b_size, G_input_dim).to(device))

# compute the loss for generator network
output_fake = D_net(data_fake)
loss_g = criterion(output_fake, real_label)

## back propagation
G_optimizer.zero_grad()
loss_g.backward()
G_optimizer.step()

## store the loss of each iter
Loss_D.append(loss_d.item())
Loss_G.append(loss_g.item())

epoch_loss_g = np.mean(Loss_G) # mean generator loss for the epoch
epoch_loss_d = np.mean(Loss_D) # mean discriminator loss for the epoch
epoch_end_time = time.time()
per_epoch_ptime = epoch_end_time - epoch_start_time

print("Epoch %d of %d with %.2f s" % (epoch + 1, epochs, per_epoch_ptime))
print("Generator loss: %.8f, Discriminator loss: %.8f" % (epoch_loss_g, epoch_loss_d))

```

```

path = image_save_dir + '/MNIST_GAN_' + str(epoch + 1) + '.png'
show_result(G_net, create_noise(25, 100).to(device), (epoch + 1), save=True, path=path)

# record the loss for every epoch
train_hist['G_losses'].append(epoch_loss_g)
train_hist['D_losses'].append(epoch_loss_d)
train_hist['per_epoch_ptimes'].append(per_epoch_ptime)


end_time = time.time()
total_ptime = end_time - start_time
train_hist['total_ptime'].append(total_ptime)

print('Avg per epoch ptime: %.2f, total %d epochs ptime: %.2f' % (
    np.mean(train_hist['per_epoch_ptimes']), epochs, total_ptime))
print("Training finish!... save training results")
with open(save_dir + '/train_hist.pkl', 'wb') as f:
    pickle.dump(train_hist, f)

show_train_hist(train_hist, save=True, path=save_dir +
'/MNIST_GAN_train_hist_epoch_chaged.png')

```

Screenshots:



Here we changes the epoch to 100

```
File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 O
+ -> Run C Markdown
1.4 Change the training epoch to 200, run it.

In [14]: # training parameters : change here
         batch_size = 100
         learning_rate = 0.0002
         epochs = 200

In [15]: # parameters for Models
         image_size = 28
         G_input_dim = 100
         G_output_dim = image_size * image_size
         D_input_dim = image_size * image_size
         D_output_dim = 1
         hidden_size = 32

In [16]: # construct the dataset and data loader
         transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize(mean=(0.5,), std=(0.5,))])
         train_data = datasets.MNIST(root=data_dir, train=True, transform=transform, download=True)
         train_loader = DataLoader(dataset=train_data, batch_size=batch_size, shuffle=True)

In [17]: # declare the generator and discriminator networks
         G_net = Generator(G_input_dim, G_output_dim).to(device)
         D_net = Discriminator(D_input_dim, D_output_dim).to(device)

         # Binary Cross Entropy Loss function
         criterion = nn.BCELoss().to(device)

         # Initialise the Optimizers
         G_optimizer = torch.optim.Adam(G_net.parameters(), lr=learning_rate)
```

End of training: took too long

```
File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 O
+ -> Run C Markdown
train_hist['per_epoch_ptimes'].append(per_epoch_ptime)

end_time = time.time()
total_ptime = end_time - start_time
train_hist['total_ptime'].append(total_ptime)

print('Avg per epoch ptime: %.2f, total %d epochs ptime: %.2f' % ( np.mean(train_hist['per_epoch_ptimes']), epochs, total_ptime))
print("Training finish!... save training results")
with open(save_dir + '/train_hist.pkl', 'wb') as f:
    pickle.dump(train_hist, f)
show_train_hist(train_hist, save=True, path=save_dir + '/MNIST_GAN_train_hist.png')

Epoch 198 of 200 with 9.50 s
Generator loss: 1.93713774, Discriminator loss: 0.80078398

Epoch 199 of 200 with 9.65 s
Generator loss: 1.93163282, Discriminator loss: 0.80723900

Epoch 200 of 200 with 11.11 s
Generator loss: 1.92899526, Discriminator loss: 0.80617598
Avg per epoch ptime: 9.76, total 200 epochs ptime: 1951.86
Training finish!... save training results
```

